

10 Putting the “Social” Back in Social Psychology

Colin Klein

1. Introduction: Two Views of Social cognition

I agree with a lot of what Victor Kumar says. Like Kumar, I agree that cognitive science can inform ethical theorizing. I think Kumar has identified several plausible routes by which cognitive science can do so. Abstracting from the details, anyone interested in the ethical life ought to care about what we *can* do and what we *typically* do when we deliberate about moral questions. The fact that our moral intuitions might be shaped by something other than the moral facts ought to trouble moral theorists (Greene, 2008). While it's possible to make that argument badly (as Berker, 2009 shows), it seems to me that Kumar has avoided the obvious pitfalls. I enthusiastically welcome Kumar's call for a move away from thinking of moral judgment as a capacity which manifests in a single discrete episode, and towards thinking of it as a dynamic, interactive, and ongoing process embedded in social contexts. Most of our important moral reasoning is done in extended conversation with others, and even solitary reasoning often embodies a dialectic structure (Stein, 2005; Laden, 2010; Campbell and Kumar, 2012).

Insofar as Kumar and I disagree, it's over how to interpret the empirical facts. Very roughly, I think that Kumar doesn't pay *enough* attention to the effects of social context. In that regard he is in good company. Insofar as we disagree, it is a symptom of a broader debate between two different models in social psychology. Kumar sides with one; I'm going to spell out the other, and show its consequences.

Here is a concrete example. Among folks of our education and class, it is vanishingly rare that one hears racist statements openly made in public life. This is (I'm told) quite a shift over a few decades back. Yet widespread racial disadvantage obviously still exists. What explains this mismatch between public avowals and private actions?

One explanation, which has captured the popular and philosophical imagination (Greenwald et al., 1998; Gendler, 2011; Strohminger et al., 2014), is that racism still exists in an *unconscious* form. It manifests in faint reactions, imperceptibly skewed judgments, the subtle recoil. The underlying cognitive states may not even be worthy of the name *beliefs*, given that they are implicit and associative (Gendler, 2008, Levy, 2014). While any particular implicit attitude may have a minuscule effect, implicit racism is so widespread and so insidious that it has profound consequences in the aggregate.

Here is a second explanation. Many people are still pretty racist. They know they are racist (or, at least, their racist beliefs are ordinary, perfectly accessible beliefs, whether or not they identify them *as* racist). They know you have to keep that sort of thing to yourself, though. People get very mad about open racism (even people who are actually racist themselves). The social costs just aren't worth it. *Quiet* discrimination persists.

These possibilities aren't strictly exclusive, but they point to two very different models of how one ought to understand social cognition. Following Fiske and Taylor (2013; also Fiske, 1993a), I'll call this the difference between thinking of humans as *Activated Actors* versus as *Motivated Tacticians*.

Activated Actors are pushovers, even despite their best intentions. As Fiske and Taylor put it, social environments "rapidly cue perceivers' social concepts, without awareness, and almost inevitably cue associated cognitions, evaluations, affect, motivation, and behavior." Hence people are hostage to "fast reactions, variously viewed as implicit, spontaneous, or automatic indicators of responses unconstrained by perceiver volition" (2013, p. 15). We may not *want* to be pushed around. But the only thing we can do is try to preempt or override our preponent reactions: Implicit responses are stubbornly resistant to top-down influence.

Motivated Tacticians, by contrast, are continually and strategically evaluating their social environment. Within a social context, agents are always "choosing among a number of possible strategies, depending on current goals" (Fiske, 1993a, p. 172). Explicit motivations thus make an important difference: Agents are constantly evaluating and re-evaluating social contexts based on their appraisals and their current goals.

One should understand these two pictures as broad poles within which different historical theories have fallen and between which social psychology has historically swung (Kihlstrom, 2004; Turiel, 2010; Fiske and Taylor, 2013). However, the difference between Activated Actors and Motivated Tacticians doesn't line up neatly with many distinctions that philosophers care about. Both emphasize the importance of unconscious processing: The Motivated Tactician also thinks that our appraisals of social situations and their demands can happen quickly and automatically (Bargh and Morsella, 2008). Both deny a classical picture of rationality on which we are objective and dispassionate agents. Within both accounts, there is disagreement about whether we are rational in some other useful sense.

Recall that both Kumar and I think that moral reasoning is a process, undertaken in an explicit or implicit social context. Moral reasoning is a species of social reasoning. Hence, on the Motivated Tactician model, we should also expect moral reasoning to be affected by social and tactical considerations. I think this possibility is both plausible and should make a difference to how we understand the empirical literature on social and moral cognition. So I will address several of the same phenomena that Kumar touches on, showing how the Motivated Tactician account gives a very different *philosophical* perspective on the same set of empirical phenomena.

Before I get down to it, a final bit about my own motivations. Activated Actor theories are interesting and worth taking seriously. But they have also had, in my opinion, the unfortunate effect of obscuring important social facts. Writing on the

research into implicit bias, for example, columnist David Brooks (2013) claimed: “Sometimes the behavioral research leads us to completely change how we think about an issue. For example, many of our anti-discrimination policies focus on finding the bad apples who are explicitly prejudiced. In fact, the serious discrimination is implicit, subtle and nearly universal.”

Many would be surprised to hear that only implicit discrimination remains. Less than six months before Brooks wrote the above, the Associated Press released a survey of the US public on various racial attitudes. Among the highlights: 37% of respondents agreed with the claim that most blacks were lazy, 38% that they were irresponsible, and 42% that they were violent. Nor is this limited to the sort of *hoi polloi* who find internet comment sections inexplicably attractive. There is ample evidence of explicit racism even among those in power.¹ Further, there appears to have been relatively little movement in racial attitudes over the past decade, even among white liberal voters (Hutchings, 2009).

Contra Brooks, there appears to be plenty of work left to be done on the explicit front. The Activated Actor model is currently over-represented when philosophers talk about social psychology. I will argue that it is time to step back and take the “social” part of social psychology more seriously.

2. World Enough and Time

Both the Motivated Tactician and Activated Actor models are *cognitive* accounts. They are methodologically individualist, focusing on internal processes.² They differ in the sort of cognitive processes they appeal to, because they differ in the parameters that show up in their explanations. Very roughly, Activated Actor models are *very* individualist: The social world is an important source of input, but once we fix the proximal input we can largely bracket off the social facts in our explanations. Motivated Tacticians, by contrast, give a much greater explanatory weight to social factors, especially variable or long-term ones.

An example will illustrate the point. Both views think that ordinary humans are cognitively limited, especially in social situations. In most circumstances, we can’t process everything with complete accuracy in real time. These limitations are a crucial part of our (well-documented) failures. Yet the *source* of those limitations differs on the two models.

Roughly speaking, the Activated Actor model continues a tradition on which observed limitations are fundamental architectural constraints of the cognitive system. Optimal behavior being impossible, we have a bias towards strategies and heuristics that rapidly get us to good enough answers (Simon, 1996). There is considerable debate *within* this tradition about whether these heuristics are innate or learned, adaptive or maladaptive, and so on (Kelman, 2013). On all versions, however, these limitations are fixed features of our cognitive apparatus.

By contrast, the Motivated Tactician picture is one on which limitations are (roughly speaking) the result of strategic interaction between numerous incompatible goals. Fiske reviews a number of different ways of cashing out these tradeoffs, describing the general problem as the tradeoff between “accuracy-oriented or open-minded motivation with confirmatory or close-minded motivation” (1993a,

p. 172). There is some sense in which (say) the tradeoff between accuracy and speed depends on our makeup (the gods need not compromise). Yet the details of these tradeoffs, Fiske emphasizes, are sensitive to what we want. As such, the parameters determining the costs of various tradeoffs are not fixed: They vary both with our appraisal of the social situation and with our motivations.³

The empirical difference is important. Consider explanations of gender stereotyping. On Activated Actor models, stereotypes are *triggered*. They enter into deliberation early, and unless they're suppressed they lead to unwanted consequences (Bargh & Chartrand, 1999). Stereotypes work like other forms of priming. Writing on stereotypes, for example, Bargh and Morsella claim that: "In contextual priming, the mere presence of certain events and people automatically activates our representations of them, and concomitantly, all of the internal information (goals, knowledge, affect) stored in those representations that is relevant to responding back" (2008, p. 76). That is an extremely general account—the same story is presumably true about non-social priming effects like (e.g.) semantic and associative priming (Stanovich & West, 1983; Seidenberg et al., 1984). Explanations of stereotypes don't actually require any special knowledge about the social milieu of the actor (other than the assumption that it was complex enough to transmit stereotypes).

Contrast this with Fiske's (1993a) Motivated Tactician account of how stereotypes enter into deliberation. Fiske notes that the prevalence and effect of stereotyping is strongly modulated by power: Powerful people tend to rely on stereotypes when dealing with subordinates, whereas subordinates stereotype the powerful far less often. She explains these facts in terms of mutually reinforcing processes of control and attention. The powerful have more demands on their time. It is comparatively more costly for them to pay attention to subordinates. Hence they must lean more heavily on fast strategies like stereotyping. Hence: "People in power stereotype in part because they do not need to pay attention, they cannot easily pay attention, and they may not be personally motivated to pay attention" (1993b, p. 621).

The opposite is true of the powerless. They have more motivation to attend to people who exert control over their lives. They have especially good reason to pay attention to the idiosyncrasies of those in power rather than treating them as mere instances of a type. Hence the asymmetry.

None of these factors, Fiske notes, need be obvious to the actors themselves. Further, stereotypes exert control not just directly, but by serving as an "implicit anchor" of which everyone in the culture is aware (1993b, p. 623). This means that stereotypes can exert power not (just) because they are widely believed, but because it is widely believed that they are widely believed. In an elegant study on the effects of beliefs about individual "brilliance" on participation in academic fields, Leslie et al. note that stereotyping may result in exclusion of women and African Americans both via overt prejudice and simply by discouraging "participation among members of groups that are currently stereotyped as not having this sort of brilliance" (2015, p. 265). That discouragement can, presumably, occur regardless of whether the stereotyped group actually *believes* the stereotype. Regardless of what you think of your own ability, it is unpleasant to contemplate

a future where you constantly have to prove yourself to people who wrongly assume that you're inadequate for the job.

Motivated Tactician models of stereotyping thus incorporate a wide variety of social facts in their explanations while remaining fully cognitive. Such accounts give correspondingly different recommendations about effective interventions. Fiske emphasizes the role of attention in modulating social expectations, for example, and the possibility that "if people pay more attention, at least some of them are less likely to stereotype" (1993b, p. 627). This also leads to differential imperatives depending on one's role in the organization: The powerful have both more need and more opportunity to change.

3. Terrible People, Terrible Reasoning

Fiske's account of stereotyping brings contingent social facts in to explain a phenomenon with obvious moral upshot. The shift to a Motivated Tactician model has a broader consequence as well: It shows the strong effects that context can have on the process of moral deliberation itself.

As I noted above, Kumar emphasizes a process that he calls "consistency reasoning" (see also Campbell and Kumar, 2012). Consistency reasoning consists in "treating like cases alike." Crucially, it is usually a *social* process that involves give-and-take with one's interlocutors. Kumar argues (plausibly, in my opinion) that consistency has been an engine for progressive social change. Yet I think consistency reasoning is far more complex in practice.

A fictional example will illustrate the point. Consider the following exchange from the television show *The Simpsons*:⁴

BART: Uh, say, are you guys crooks?

FAT TONY: Bart, is it wrong to steal a loaf of bread to feed your starving family?

BART: No.

TONY: Well, suppose you got a large starving family. Is it wrong to steal a truckload of bread to feed them?

BART: Uh uh.

TONY: And, what if your family don't like bread? They like ... cigarettes?

BART: I guess that's okay.

TONY: Now, what if instead of giving them away, you sold them at a price that was practically giving them away. Would that be a crime, Bart?

BART: Hell no!

This seems to me to be a clear case of moral consistency reasoning. It is also a *failed* case. Bart has gotten to the wrong result in the wrong way. But why?

The Activated Actor model might cite ways in which Bart was *pushed* away from the obvious truth. Perhaps he relied on unreliable moral heuristics (Sunstein, 2005). Perhaps he put too much weight on the output of fast unconscious processing (Kahneman, 2011). Perhaps he was easily swayed by the emotional pull of Fat Tony's argument (Greene, 2008). Perhaps he fell prey to a morally troubling

in-group bias (Gino and Galinsky, 2012) or some other set of evolutionarily based biases (Wielenberg, 2010). Perhaps all these and more. But of course, then the knife twists: *We do the same thing*. Bart’s failure differs from ours only in degree, not in kind. Even that is cold comfort, for most of these biases are hard to recognize from the inside. Bart’s final confidence is as characteristic as his failures.

Yet it seems to me that this can’t be all that’s gone wrong. Bart also made a more basic mistake: He assumed that Fat Tony was *arguing in good faith*. It may be reasonable to extend that sort of charity to our interlocutors, at least initially. But it’s also an assumption that we should be willing to abandon—especially if we hear a tortured set of analogies that leads to a suspiciously self-serving end. In such cases, sensitive moral agents should be inclined to rescind their initial charity, and stick to their initial judgments.

Moral deliberation is difficult precisely because we live in a world with a lot of bad people. Many people are self-serving liars, at least some of the time. Many people care more about making peace than about getting to the right moral answer. Most moral consistency reasoning also isn’t done in the seminar room; it is done when one person has harmed another, and is trying to get off the hook. Sometimes people really care about the moral truth. Very often, they care about avoiding consequences.

That makes moral consistency reasoning very easy to *exploit*. This goes both ways. Most confessions to the police (at least from the actually guilty) are not the result of coercion. Interrogators rarely need the phone book under the hot lights. Instead, as Zulawski and Wicklander (2002) point out, the most effective interrogators simply help the accused *rationalize* what they did (“He treated you like a slave rather than a valued employee. The business was going under anyway, and no one would miss the money. So it’s not like you hurt anyone, and you were only trying to pay the mortgage ...”). The drive to find a narrative on which your actions come out permissible is a very strong one.

Or consider the incredible vitriol that hypocrisy attracts. At first pass, it’s not clear why hypocrisy should be so troubling. There is a mismatch between the hypocrite’s explicit moral judgments and their behavior. But as Wallace (2010) notes, mere inconsistency isn’t normally treated as a terrible personal failing. Nor does it seem like personal inconsistency ought to undermine first-order moral judgments. Yet people often take hypocrisy as evidence of faulty reasoning. Why? The answer, I suggest, is that judgments of hypocrisy play a crucial functional role in regulating factors that would otherwise tend to undermine moral discourse. Wallace (2010) argues that hypocrisy judgments provide a check on inappropriate instances of the reactive attitudes: We blame the hypocrite for their own unjustified blame. Isserow and Klein (2017) further argue that the role of hypocrisy judgments is to undermine putative claims to *moral authority*. We often defer to others’ moral judgments, particularly to those we think we can trust. But that puts one in a morally precarious situation, open to exploitation. Sensitivity to hypocrisy is necessary to regulate that vulnerability both for ourselves and for others.

Note, again, that while these are all cognitive facts, they depend on contingent social facts. It is only when we take into account the motivations of others that

we have to start worrying about lying and manipulation. It is only when we place some of the epistemic costs of moral deliberation onto perceived authorities that we have to start looking for hypocrites. And it is only when we try to accommodate the actual social facts about differential power structures that we have to think about potential asymmetries in attention, access, and control.

Does this show that Kumar's moral consistency reasoning is wrong, or limited in scope? I think not. Rather, it shows that it operates at two different levels. There is a global level, on which we want our moral principles and particular judgments to be in agreement. That is the sense of moral consistency reasoning we have in mind when we consider large-scale changes in our beliefs.

But consistency reasoning, I submit, also operates in *local* contexts. That is, we also want our beliefs about the local social situation to cohere. Much of our social and moral discourse is aimed at those local contexts, not at big-picture moral considerations.

The way we frame the situation, the judgments we make, and the justifications we offer to others should, ideally, come into equilibrium with our beliefs about the people we're with, the sort of arguments we expect them to find convincing, and the goals we think that moral conversation is trying to achieve. Sometimes the correct response to perceived inconsistency is to change our judgments or interpretation of a difficult case. Other times, we should change our beliefs about what our interlocutors think, or what the point of arguing over a moral point really is.

4. Experimental Situations Are Social Situations

Local consistency reasoning is interesting in its own right. It is also interesting because it gives the Motivated Tactician account an alternative take on many experiments in the cognitive science of moral judgment. As an example, consider Jonathan Haidt's well-known work on moral dumbfounding. Haidt and his collaborators constructed cases like the following:

Julie and Mark, who are brother and sister, are traveling together in France. They are both on summer vacation from college. One night they are staying alone in a cabin near the beach. They decide that it would be interesting and fun if they tried making love. At very least it would be a new experience for each of them. Julie was already taking birth control pills, but Mark uses a condom too, just to be safe. They both enjoy it, but they decide not to do it again. They keep that night as a special secret between them, which makes them feel even closer to each other. So what do you think about this? Was it wrong for them to have sex?

(Haidt et al., 2000, p. 20)

They find that subjects are often "dumbfounded": They judge that the case is wrong, they are unable to offer any good reasons for thinking so, and yet they stubbornly persist in their judgment (Haidt et al., 2000). This seems to be good evidence for a non-rational, non-deliberative, social intuitionist view of moral judgment (Haidt, 2001). While there are problems with the original (unpublished)

study, I take it that the *phenomenon* is familiar enough: We’ve all met people who simply dig in their heels and stick to their guns on difficult moral issues.

The real question is *why*. Here I think there is a further benefit to thinking about humans as Motivated Tacticians. For starters, we are deeply concerned not just with *what* people do but why they do it. That is especially important when we turn to imagining morally laden actions: As Kennett (2011) puts it, “In order for us to imaginatively enter another’s point of view, we must be able to see the *point* of what they do” (2011, p. 182). An inability to do so engenders imaginative resistance. But as Kennett points out, the only reason why Julie and Mark want to sleep together is that it would be a “new experience” (2011, p. 185)—a difficult motivation to accept in the case of such a loaded taboo.

Indeed, Royzman et al. (2015), revisiting Haidt’s study, found that most subjects simply didn’t believe that Mark and Julie’s actions would not be harmful. Liu and Ditto (2013) similarly found that moral evaluation had an effect on factual appraisal of dilemmas: Subjects were less willing to accept parts of the setup of moral dilemmas that conflicted with their overall moral judgment. Put another way: Even explicit instruction is one more, *non-privileged*, input into local consistency reasoning. It can be overridden if it’s too implausible.

I think there is more to say. Given the right motivations, there can be a disconnect between the reasons you *have* and the reasons you *give*. That is true in any social situation. But there is good evidence that experimental situations—particularly ones that involve dialogue with experimenters—are social situations with unique demands. Milgram’s (1963) work on experimental obedience is well-known, of course, but the effect can be more subtle. As Orne (1962) demonstrated, subjects will eagerly interpret experimental demands in a way that makes them meaningful, and will do so even if they are explicitly informed that they are engaged in a pointless task.

Further, experimental contexts are also *communicative* contexts. As such, subjects interpret and give information in conformance with Gricean conversational norms. Subjects tend to assume, in accordance with Gricean norms, that all of the information they are given is relevant.

This is a well-known confound in studies of heuristics and biases (Hilton, 1995; Kelman, 2013). Apparently irrational behavior can often be redeemed if we pay attention to conversational pragmatics. Suppose I ask my wife whether (hypothetically speaking) she would be mad if I slept with one of her friends. That might make her a little suspicious about my fidelity. Suppose I asked her whether (hypothetically speaking) she would be mad if I slept with her friend Natalie. She would likely become *very* suspicious. Ought we convict my wife of irrationality? Given that Natalie is one of her friends, the latter can’t possibly be more likely than the former. But of course, the fact that I specifically mentioned *Natalie* demands an explanation—conversational norms require us to be as specific as necessary and no more. The most obvious explanation is that I have Natalie on my mind. Attention to conversational pragmatics thus redeems what looks like a simple failure to appreciate probabilities.

The same processes occur in empirical studies of moral reasoning. Subjects tend to give only information that they think will be relevant to the experimenter. In

a nice study of this effect, Norenzayan and Schwartz (1999) looked at subjects' causal attributions about the actions of a mass murderer. They found that subjects gave more situational explanations when told that they were talking with a social scientist, and more dispositional explanations when they thought they were talking to a personality psychologist.

In sum, the fact of being in an experimental situation can have non-negligible influences on how subjects behave. That effect, I argue, ultimately stems from the effect of local consistency reasoning. Subjects want their reaction to particular cases to be consistent with both their general beliefs about the world and their particular beliefs about the situation they are in.

With that in mind, here's an alternative interpretation of Haidt's experiment. Subjects faced with the Julie and Mark case are asked whether that action was wrong. They judge that it's incest, and incest is wrong. That's all pretty obvious; the experimenter has foregrounded the incest-relevant bits. Gricean norms thus make the stated question puzzling. Nor does it seem like they are being asked to justify the *principle* that incest is always wrong, as opposed to passing judgment on a particular case. The experimenter is a member of the subjects' moral community, and they have no special reason to doubt that the experimenter also thinks that incest is wrong. Nor does the experimenter seem seriously invested in changing the subjects' moral beliefs on the topic. Hence subjects can rely on what Tetlock (1992) calls the "acceptability heuristic," according to which least-effort solutions to a moral problem are probably the correct ones. Given this, subjects focus less on whether the action was *wrong* and more on whether Julie and Mark ought to be *blamed*—that is, on the question about their willingness to enforce judgments in complicated cases.

Once you get to that point, it might be quite reasonable to get hard-headed and to ignore putative excuses. In a discussion of the fundamental attribution error, Tetlock notes:

One way of pressuring other people to behave is by indicating to them that one has a low tolerance for justifications or excuses and that one will treat their behavior as automatically diagnostic of underlying intentions and personality attributes ... As Axelrod notes, people are not only expected to act in accord with prevailing norms, they are also expected to censure those who violate norms (a norm to enforce norms, or a metanorm). Insofar as accountable subjects feel that their moral mettle is being tested (their willingness to apply metanorms), they may be more motivated to hold others responsible and to reject situational explanations or excuses ... The argument simply asserts that accountable subjects will tend to rely on the acceptability heuristic when they can infer the views of the prospective audience. In this case, subjects may assume that the prospective audience expects them to hold people responsible for deviant or untoward conduct.

(Tetlock, 1992, p. 361)

It is not unusual that Julie and Mark have a good excuse for violating a moral norm. *Everyone* has an excuse. Good moral agents don't take those excuses at face

value, even if experimenters would like them to. In Royzman et al.’s (2015) study on dumbfounding, they found that the instance of true dumbfounding dropped to virtually zero when subjects were carefully asked about their commitments and which question they thought they were answering.

5. Moral Development

This leads to a final point. Explanations like Fiske’s of stereotyping or mine of dumbfounding cite differential power relationships *between* groups. But of course there are also likely to be substantial and morally important differences *within* groups: differences of motivation, of personality, of inclination, and of ability. The Motivated Tactician account paints moral deliberation—particularly in social contexts—as relying on forms of social competence. People obviously differ in their social competence. One might wonder whether those differences also make a difference to judgments.

Activated Actor accounts have largely downplayed individual differences. They belong to an experimental tradition that seeks, in Bakan’s (1967) terms, *general* facts that hold of all subjects, not just averages about aggregates. At best, individual differences are explained in terms of variation of these basic elements. Usually, “individual variation is a source of embarrassment to the experimenter,” mere error variance that obscures the main effect of interest (Cronbach, 1957, p. 674). As Cronbach (1957) points out, there is a second tradition in psychology that is concerned with correlation and explanation of individual differences.⁵

One obvious place that individual differences become important is in the study of moral development. Following theorists like Piaget (1932) and Kohlberg (1969), studies of moral judgment have assumed that moral judgment proceeds in distinct stages rather than springing forth fully formed. Further, moral competence might be the sort of thing that takes time and effort to develop. In a famous passage from the *Analects* about his own moral development Confucius claimed that:

At fifteen, I had my mind bent on learning. At thirty, I stood firm. At forty, I had no doubts. At fifty, I knew the decrees of Heaven. At sixty, my ear was an obedient organ for the reception of truth. At seventy, I could follow what my heart desired, without transgressing what was right.⁶

Moral development is a lifelong process, so one might reasonably expect people to be much better (or worse) at it at different points in their lives.

Further, as recent research has emphasized, development is a process that also occurs in dialogue between children and their peers, their parents, and other authorities. This dialogue, like other forms of cultural learning, occurs in carefully scaffolded environments (Sterelny, 2012). It also opens the possibility that children are far from passive recipients of moral dictates. Instead, as Turiel emphasizes:

From the beginning, the developing person acts to give meaning to and to understand the world, and to make sense of social relationships. And with development, they come to evaluate the dictates of others, to distinguish

between what they judge to be legitimate and illegitimate demands and directives, and to engage in relational reciprocal interpersonal interactions and communications. In the process of early development, they evaluate social relationships and consequently come to accept and reject parental directives, sometimes leading to conflicts with adults and peers. Children, like adults, also approach social relationships from the viewpoint of moral concerns, including concerns with the welfare of others and fairness.

(2010, p. 559)

This long developmental process is thus a deeply *social* one, reflecting a lifelong engagement in both global and local consistency reasoning.

As such, the process of moral development—and our particular judgments in the short term—are likely to be deeply shaped by long-term facets of that society. There is an increasing body of research on the content and structure of this shaping process (Graham et al., 2011). This includes the effect of political (Graham et al., 2009) and religious (Hofmann et al., 2014) beliefs, and the broader cultural and socioeconomic background (Henrich et al., 2010).

Some have taken this variation to be an argument against the stronger forms of moral realism (Machery et al., 2005), citing precisely the sorts of considerations that Kumar addresses. Thinking of humans as Motivated Tacticians isn't a particularly happy place for the moral realist to rest: If moral reasoning is sensitive to tactical and strategic considerations, even our considered philosophical judgments may not rise above our time, place, and role.

I'd like to conclude on a slightly more optimistic note—one that I hope will be compatible with the points on which Kumar and I both agree. If we think of humans as Activated Actors, then variation is indeed problematic: People either get it right or they don't, and there's little more to say. If moral reasoning is social reasoning, however—if, as Fiske and Taylor put it, "[people's] social thinking is for their social doing" (2013, p. 15)—then the mere presence of variation doesn't mean that everyone is getting things wrong.

Social life is complex, motivations are diverse, and resources are widely distributed. We should recognize the profound influence that the social has on the personal, and from there on the moral. But once we do we can also think about how we might *construct* social contexts that are less susceptible to the most damaging of social effects. I've argued that some paradigmatic social psychology experiments failed to do so. But philosophers and cognitive scientists are well-poised to work out the intermingled effects of the social on the cognitive. What is required is turning our gaze outward, towards the ways in which social factors affect our moral judgments.⁷

Notes

- 1 As I write, the state of Pennsylvania is dealing with revelations that high-profile state officials, including judges and state prosecutors, regularly traded racist and sexist jokes over official email (*Washington Post*, Dec. 26, 2015, "Pornographic email scandal roils Pennsylvania politics"). There is also considerable debate about whether implicit biases

- as measured by the IAT (Implicit Association Test) actually correlate with negative outcomes in a way that’s not simply mediated by explicit bias; see Oswald et al. (2013) for a review.
- 2 There is also a rich tradition of methodologically holistic, social-functionalist approaches to morality—e.g., those drawing on Durkheim (1915) or the Marxist corpus. These have had far less influence on analytic philosophy, at least in recent years.
 - 3 In this, the model fits well with more recent approaches to automaticity, which have stressed a continuous and strategic boundary between the automatic and the controlled (Bargh and Morsella, 2008; Kihlstrom, 2008). Further afield, there is an interesting parallel to recent work on cognitive control, which stresses that capacity limitations arise from the tradeoffs between efficiency and effectiveness in the control system, rather than being primitive architectural features (Botvinick and Cohen, 2014).
 - 4 Episode 8F03, “Bart the Murderer.” For those unfamiliar with the series: Bart is a young larrikin who has fallen in with mafia boss Fat Tony and his gang. Asked to hide a hijacked shipment of cigarettes, he has begun to worry about his new friends’ motivations.
 - 5 Cronbach associates this tradition most strongly with research on personality; while this is historically plausible, it is worth noting that there are personality researchers who focus on common individual processes (cf. Cervone, 2005). Cronbach also makes the fascinating suggestion that individual differences research—growing as it did out of research on IQ and industrial placement—tends towards the politically conservative, while the search for common features is liberal and democratic.
 - 6 *Analects* Book 2, Chapter 4. Legge translation.
 - 7 Thanks to Jessica Isserow and Victor Kumar for helpful comments on this draft, as well as audiences at the 2011 Pacific APA and the Social Psychology Brownbag Series at the University of Illinois at Chicago for helpful feedback on a distant ancestor.

References

- Bakan, D. 1967. The General and the Aggregate. In his *On Method: Toward a Reconstruction of Psychological Investigation*, 34–6. San Francisco, CA: Jossey-Bass, Inc.
- Bargh, J. A., & Chartrand, T. L. 1999. The unbearable automaticity of being. *American Psychologist* 54: 462–79.
- Bargh, J. A., & Morsella, E. 2008. The unconscious mind. *Perspectives on psychological science* 3: 73–9.
- Berker, S. 2009. The Normative Insignificance of Neuroscience. *Philosophy & Public Affairs* 37: 293–329.
- Botvinick, M. M., & Cohen, J. D. 2014. The computational and neural basis of cognitive control: Charted territory and new frontiers. *Cognitive Science* 38: 1249–85.
- Brooks, D. 2013. Beware Stubby Glasses. *New York Times*, Jan. 10, 2013.
- Campbell, R., & Kumar, V. 2012. Moral Reasoning on the Ground. *Ethics* 122: 273–312.
- Cervone, D. 2005. Personality architecture: Within-person structures and processes. *Annual Review of Psychology* 56: 423–52.
- Cronbach, L. J. 1957. The two disciplines of scientific psychology. *American Psychologist* 12: 671–84.
- Durkheim, E. 1915. *The Elementary Forms of Religious Life*. New York: Macmillan.
- Fiske, S. T. 1993a. Social cognition and social perception. *Annual Review of Psychology* 44: 155–94.
- Fiske, S. T. 1993b. Controlling other people: The impact of power on stereotyping. *American Psychologist* 48(6): 621–8.

- Fiske, S. T., & Taylor, S. E. 2013. *Social Cognition: From Brains to Culture*. London: Sage Press.
- Gendler, T. S. 2008. Alief and Belief. *Journal of Philosophy* 105: 634–63.
- Gendler, T. 2011. On the epistemic costs of implicit bias. *Philosophical Studies* 156: 33–63.
- Gino, F., & Galinsky, A. D. 2012. Vicarious dishonesty: When psychological closeness creates distance from one's moral compass. *Organizational Behavior and Human Decision Processes* 119: 15–26.
- Graham, J., Haidt, J., & Nosek, B. A. 2009. Liberals and conservatives rely on different sets of moral foundations. *Journal of Personality and Social Psychology* 96: 1029–46.
- Graham, J., Nosek, B. A., Haidt, J., Iyer, R., Koleva, S., & Ditto, P. H. 2011. Mapping the moral domain. *Journal of Personality and Social Psychology* 101: 366–85.
- Greene, J. D. 2008. The secret joke of Kant's soul. In W. Sinnott-Armstrong (Ed.), *Moral Psychology*, 35–80. Cambridge, MA: MIT Press.
- Greenwald, A. G., McGhee, D. E., & Schwartz, J. L. K. 1998. Measuring individual differences in implicit cognition: The implicit association test. *Journal of Personality and Social Psychology* 74: 1464–80.
- Haidt, J. 2001. The emotional dog and its rational tail. *Psychological Review* 108: 814–34.
- Haidt, J., Bjorklund, F. & Murphy, S. 2000. Moral dumbfounding: When intuition finds no reason. Unpublished manuscript, University of Virginia.
- Henrich, J., Heine, S. J., & Norenzayan, A. 2010. The weirdest people in the world? *Behavioral and Brain Sciences* 33: 61–83.
- Hilton, D. J. 1995. The social context of reasoning: Conversational inference and rational judgment. *Psychological Bulletin* 118: 248–8.
- Hofmann, W., Wisneski, D. C., Brandt, M. J., & Skitka, L. J. 2014. Morality in everyday life. *Science* 345: 1340–3.
- Hutchings, V. L. 2009. Change or more of the same? Evaluating racial attitudes in the Obama era. *Public Opinion Quarterly* 73: 917–42.
- Isserow, J., & Klein, C. 2017. Hypocrisy and Moral Authority. *Journal of Ethics and Social Philosophy* 12: 191–222.
- Kahneman, D. 2011. *Thinking, Fast and Slow*. New York: Macmillan.
- Kelman, M. 2013. Moral realism and the heuristics debate. *Journal of Legal Analysis* 5: 339–97.
- Kennett, J. 2011. Imagining reasons. *Southern Journal of Philosophy* 49: 181–92.
- Kihlstrom, J. F. 2004. Is there a “People are Stupid” school in social psychology? *Behavioral and Brain Sciences* 27: 348.
- Kihlstrom, J. F. 2008. The automaticity juggernaut—or, are we automatons after all? In J. C. Kaufman, R. F. Baumeister, & J. Baer (Eds.), *Are We Free? Psychology and Free Will*, 155–80. New York: Oxford University Press.
- Kohlberg, L. 1969. Stage and sequence: The cognitive-developmental approach to socialization. In D. Goslin (Ed.), *Handbook of Socialization Theory and Research*, 110–20. Chicago, IL: Rand McNally.
- Laden, A. 2010. Reasoning: It's not all in the head. *Human Development* 53: 105–9.
- Leslie, S.-J., Cimpian, A., Meyer, M., & Freeland, E. 2015. Expectations of brilliance underlie gender distributions across academic disciplines. *Science* 347: 262–5.
- Levy, N. 2014. Neither fish nor fowl: Implicit attitudes as patchy endorsements. *Notes* 49: 800–23.
- Liu, B. S., & Ditto, P. H. 2013. What Dilemma? Moral evaluation shapes factual belief. *Social Psychological and Personality Science* 4: 316–23.
- Machery, E., Kelly, D., & Stich, S. P. 2005. Moral realism and cross-cultural normative diversity. *Behavioral and Brain Sciences* 28: 830.
- Milgram, S. 1963. Behavioral study of obedience. *Journal of Abnormal and Social Psychology* 67(4): 371–8.

- Norenzayan, A., & Schwarz, N. 1999. Telling what they want to know: Participants tailor causal attributions to researchers' interests. *European Journal of Social Psychology* 29: 1011–20.
- Orne, M. T. 1962. On the social psychology of the psychological experiment: With particular reference to demand characteristics and their implications. *American Psychologist* 17: 776–83.
- Oswald, F. L., Mitchell, G., Blanton, H., Jaccard, J. & Tetlock, P. E. (2013). Predicting ethnic and racial discrimination: A meta-analysis of IAT criterion studies. *Journal of Personality and Social Psychology* 105(2): 171.
- Piaget, J. 1932. *The Moral Judgement of the Child*. London: Routledge & Kegan Paul.
- Racial attitudes survey. 2012, October 29. The Associated Press. http://surveys.ap.org/data%5CGfK%5CAP_Racial_Attitudes_Topline_09182012.pdf
- Royzman, E. B., Kim, K., Leeman, R. F. et al. 2015. The curious tale of Julie and Mark: Unraveling the moral dumbfounding effect. *Judgment and Decision Making* 10: 296–313.
- Seidenberg, M. S., Waters, G. S., Sanders, M., & Langer, P. 1984. Pre- and postlexical loci of contextual effects on word recognition. *Memory & Cognition* 12: 315–28.
- Simon, H. A. 1996. *The Sciences of the Artificial*. Cambridge, MA: MIT Press.
- Stanovich, K. E., & West, R. F. 1983. On priming by a sentence context. *Journal of Experimental Psychology: General* 112(1): 1–36. doi: <https://doi.org/10.1037/0096-3445.112.1.1>
- Stein, E. 2005. Wide reflective equilibrium as an answer to an objection to moral heuristics. *Behavioral and Brain Sciences* 28: 561–2.
- Sterelny, K. 2012. *The Evolved Apprentice*. Cambridge, MA: MIT Press.
- Strohminger, N., Caldwell, B., Cameron, D., Borg, J. S., & Sinnott-Armstrong, W. 2014. Implicit morality: A methodological survey. In C. Luetge, H. Rusch, & M. Uhl (Eds.), *Experimental Ethics: Towards an Empirical Moral Philosophy*, 133–56. London: Palgrave Macmillan.
- Sunstein, C. R. 2005. Moral heuristics. *Behavioral and Brain Sciences* 28: 531–41.
- Tetlock, P. E. 1992. The impact of accountability on judgment and choice: Toward a social contingency model. *Advances in Experimental Social Psychology* 25: 331–76.
- Turiel, E. 2010. The development of morality: Reasoning, emotions, and resistance. In R. M. Lerner, M. E. Lamb, & A. M. Freund (Eds.), *The Handbook of Life-Span Development*, 554–83. New York: John Wiley and Sons.
- Wallace, R. J. 2010. Hypocrisy, moral address, and the equal standing of persons. *Philosophy & Public Affairs* 38: 307–41.
- Wielenberg, E. J. 2010. On the evolutionary debunking of morality. *Ethics* 120: 441–64.
- Zulawski, D. E., & Wicklander, D. E. 2002. *Practical Aspects of Interview and Interrogation*. New York: CRC Press.

Further Readings for Part V

Demaree-Cotton, J., & Kahane, G. (2019). The neuroscience of moral judgment. In A. Zimmerman, K. Jones, & M. Timmons (Eds.), *The Routledge Handbook of Moral Epistemology*, 84–149. New York: Routledge.

Argues that neuroscience is typically relevant to moral epistemology only when it bears on cognitive science claims pitched at higher levels of descriptions and evaluates extant attempts to draw conclusions for moral epistemology on the basis of theories in the cognitive science of morality.

Gigerenzer, G. (2010). Moral satisficing: Rethinking moral behavior as bounded rationality. *Topics in Cognitive Science*, 2(3), 528–54. <https://doi.org/10.1111/j.1756-8765.2010.01094.x>

Attempts to derive conclusions about ethics from a theory on which moral behavior is shaped by social heuristics tailored to the social environment.

Kauppinen, A. (2014). Ethics and empirical psychology – critical remarks to empirically informed ethics. In M. Christen, C. van Schaik, J. Fischer, M. Huppenbauer, & C. Tanner (Eds.), *Empirically Informed Ethics: Morality between Facts and Norms*, 279–305. https://doi.org/10.1007/978-3-319-01369-5_16

Provides a pessimistic assessment of attempts to derive metaethical and ethical conclusions from empirical work.

Nichols, S. (2014). Process debunking and ethics. *Ethics*, 124(4), 727–49. <https://doi.org/10.1086/675877>

Provides an account of the conditions under which empirically based debunking arguments in metaethics and ethics are likely to be successful or unsuccessful.

Plakias, A. (2016). Metaethics: Traditional and empirical approaches. In J. Systemsma & W. Buckwalter (Eds.), *A Companion to Experimental Philosophy*. doi: 10.1002/9781118661666.ch13

Reviews attempts to use empirical work in moral psychology to make progress on debates in metaethics regarding moral realism, cognitivism, and moral motivation.

Pözlner, T. (2018). Can the empirical sciences contribute to the moral realism/anti-realism debate? *Synthese*, 195(11), 4907–30. <https://doi.org/10.1007/s11229-017-1434-8>

Provides an optimistic assessment of attempts to use empirical work to adjudicate debates in metaethics, including debates between cognitivism and non-cognitivism, and provides an account of the conditions that must be met in order for empirical work to do so.

Prinz, J. (2015). *An Empirical Case for Motivational Internalism*. Retrieved from <https://www.oxfordscholarship.com/view/10.1093/acprof:oso/9780199367955.001.0001/acprof-9780199367955-chapter-4>

Provides several empirically based arguments in favor of motivational internalism about moral judgment.

Rini, R. A. (2015). Psychology and the aims of normative ethics. In J. Clausen & N. Levy (Eds.), *Handbook of Neuroethics*, 149–68. https://doi.org/10.1007/978-94-007-4707-4_162

Defends the relevance of empirical work to adjudicating debates on normative ethics, the nature of moral agency, and the extent to which people are justified in holding their moral beliefs.

Tiberius, V. (2006). Well-being: Psychological research for philosophers. *Philosophy Compass*, 1(5), 493–505.

Discusses interactions between the philosophical and psychological study of wellbeing, while considering the relevance of this research for public policy.

Wright, J. C. (2018). The fact and function of meta-ethical pluralism: Exploring the evidence. *Oxford Studies in Experimental Philosophy*, 2, 119–50.

Argues that empirical work shows that most people are both realists and anti-realists about morality, and that this has certain benefits for society.